





# Voxels



Nick Horton







# The Humble Cube













| Data Name | Type     | Size (bytes) |
|-----------|----------|--------------|
| position  | float[3] | 12           |
| color     | float[3] | 12           |
| total     |          | 24           |

24 bytes per vertex  
3 vertex per triangle  
2 triangles per square  
6 squares per cube  
230,503 cubes per model  
16 models per scene

$32 \times 3 \times 2 \times 6 \times 230,503 \times 16$   
= ~3 billion bytes aka 3 gigabytes of geometry data.



## The Scene



On my (bad) laptop this renders at ~6 FPS.

How can we optimize this (memory and FPS)?

# Topics

- Chunking
- Memory
- Restrictions
- Instancing
- Culling
- Greedy Meshing

# Chunking

- Separation of concerns
- Finite remeshing size
- Makes frustum culling possible



**FPS: 6.6**  
**Mem: 2.9G**  
**DC: 1**



**FPS: 6.6**  
**Mem: 2.9G**  
**DC: 16**

# Memory

- (most) voxels are grid aligned
- Depending on chunk size you need  $\log_2(\text{chsz})$  bits
- Uniform for chunk offset
- For our case chunking at 256 means 1 uint can contain x,y,z with 8 bits to spare
- Same with color

**FPS: 6.6**  
**Mem: 2.9G**  
**DC: 16**



**FPS: 6.6**  
**Mem: 0.9G**  
**DC: 16**

# Restrictions

- Most game optimizations make concessions for performance
- Often these restrictions are placed on the artists/designers
- For example (in our case) using a color palette instead of rgb color for each voxel
- Index into palette requires 8 bits luckily, we have an extra byte in the position uint to store this.

**FPS: 6.6**  
**Mem: 0.9G**  
**DC: 16**



**FPS: 6.7**  
**Mem: 0.5G**  
**DC: 16**

# Instancing

- We only have 6 (4) vertexes of real geometry data, that of a square.
- Everything else is per square or per cube
  - Direction of square
  - pos and col of cube
- We have A LOT of duplicate data


**FPS: 6.6**  
**Mem: 0.5G**  
**DC: 16**



**FPS: 8.7**  
**Mem: 14M**  
**DC: 16**

# Face Culling

- Most faces we render will never be seen.
- Removing these faces is an expensive calculation but chunking makes it possible in realtime.



**FPS: 8.7**  
**Mem: 14M**  
**DC: 16**



**FPS: 43.7**  
**Mem: 14M**  
**DC: 16**

# Greedy Meshing

- Extremely expensive calculation
- Mathy to implement
- So so so fast for static meshes

